

datagraph

AI-powered Change Impact Graph for data and code systems

Complete project documentation — from idea to open-source-ready library

Version 0.5.0

Location: C:\Users\Hp\Documents\Innovation\datagraph

Built with Claude Code (Claude Fable 5) · August 2026

License: MIT · Python $\geq$ 3.9 · pip install datagraph

Contents

1. Executive summary	4
2. The brief and the design principle	4
3. How it was built — the A-to-Z timeline	5
3.2 The v0.3 build — lineage, schema relationships, and “build everything”	7
3.1 The v0.2 build — closing the gaps against Graphify and DataHub	7
4. Skills, tools and technologies used	8
4.1 Claude Code skills and tools	8
4.2 Python libraries and standards	9
5. Architecture	10
5.1 Pipeline	10
5.2 Package layout	10
5.3 Data model	11
5.4 Node id conventions	11
5.5 Risk scoring and test recommendations	11
5.6 The AI layer	12
6. Worked example — the demo	13
7. How to use it	13
7.1 Install	13
7.2 Command line	13
7.3 Python API	14
8. Testing — what was verified and how	15
8.1 Automated suite (108 tests, all passing)	15
8.2 End-to-end verification from the built wheel	15
8.3 Bugs found and fixed along the way	16
8.4 Honest remaining gaps	16
9. How datagraph differs from what already exists	18
9.1 Side by side: Graphify · DataHub · OpenLineage · datagraph	18
9.2 What Graphify has that datagraph does not	19
9.3 What DataHub has that datagraph does not	20
9.4 What datagraph has that neither does	20
9.5 What datagraph deliberately does not try to be	20
9.6 Where existing tools are ahead (known gaps)	20
10. Packaging and open-source readiness	20
10.1 Packaging	20
10.2 Repository hygiene	21
10.3 Publishing to GitHub (next step)	21

11. Roadmap..... 21

12. v0.5 — data profiling, knowledge base for AI assistants, extractor plugins 22

12.1 What was added..... 22

12.2 Why this matters..... 23

12.3 Tests added..... 23

Appendix A — Command reference 24

Appendix B — File inventory 24

Appendix C — Design decisions worth remembering..... 24

(In Microsoft Word: right-click the table of contents and choose Update Field if page numbers are missing.)

1. Executive summary

Name note: the library was developed as “impactgraph” and renamed to “datagraph” in v0.4.0 (PyPI name datagraph; GitHub sumit-gupta03/datagraph). Node ids, commands and the API are unchanged apart from the package/command name; the graph class keeps the name ImpactGraph with DataGraph as the preferred alias.

datagraph is a pip-installable Python library and command-line tool that answers the question every data-platform and backend engineer asks before merging a change:

The question

“If I change this file, function, dbt model, SQL column, API, or table — what can break?”

It builds one unified dependency graph that spans both the code world (Python files, functions, imports, calls) and the data world (dbt models, sources, warehouse tables, columns, dashboards), and it does so deterministically — from real engineering artifacts, never from an LLM’s guess. Graph algorithms then compute the blast radius of a change, a risk level, and a concrete test plan. An optional AI layer explains the computed result in plain language for the engineer about to merge.

Headline facts

- **Name:** datagraph — verified free on PyPI before building (impactlens was also free).
- **Scope delivered (v0.2):** 7 deterministic extractors (Python AST, dbt manifest.json incl. compiled-SQL column lineage, SQL via sqlglot with column-level lineage, git diff, OpenLineage events, DataHub lineage files, warehouse information_schema), unified graph with direction-aware impact propagation and edge provenance (extracted / inferred), risk scoring, owners to notify, test recommender, rich terminal report with ASCII fallback, interactive HTML view, 13 CLI commands (build, impact, diff, nodes, paths, hotspots, graph-diff, export, html, watch, hook-install, explain, mcp), MCP server, a /datagraph skill for AI coding assistants, a GitHub Action, and the optional Claude-powered explanation layer.
- **Quality:** 108 automated tests passing (including dbt’s real jaffle_shop project as a fixture); end-to-end verified from the built wheel in a clean virtual environment; one real Windows bug (UTF-8 BOM) found and fixed during that pass.
- **Open-source readiness:** MIT license, README, CONTRIBUTING guide, .gitignore, GitHub Actions CI (Ubuntu + Windows × Python 3.9/3.11/3.13), wheel + sdist built, two commits on a local main branch, ready to push.

2. The brief and the design principle

The project started from a product idea: existing tools already cover pieces of the problem — Graphify and Code-Graph-RAG build graphs from source code, while DataHub and OpenLineage focus heavily on data lineage — but nothing connects both worlds into a single “what breaks if I change this” answer. The target experience was a chain such as:

```
Git change → Python function → Lambda → API → S3 → Snowflake table → dbt model → report / dashboard
```

and a report such as:

```
⚠ Change Impact

Changed:  models/customer.sql  (customer_id → customer_key)

Potential impact:
  customer.sql → dim_customer → fact_booking → { revenue_report,
customer_dashboard }

Risk: HIGH
Affected: 3 dbt models, 2 Snowflake tables, 1 Python function, 2 reports

Recommended tests:
  ✓ dbt relationship test    ✓ fact_booking integration test    ✓ customer API
contract test
```

The one non-negotiable design principle

The LLM never creates the fundamental graph. The graph is built deterministically from actual engineering artifacts — Python AST, SQLGlot, dbt manifest.json, git diff, (later) Snowflake metadata and Airflow/Lambda definitions — then graph algorithms compute impact, and only then does AI produce the explanation, risk narrative, and test plan.

This gives far higher trust than asking a model to guess relationships, and it is the line that separates datagraph from LLM-first code-graph tools.

3. How it was built — the A-to-Z timeline

The library was built in a single autonomous Claude Code session. Below is the actual sequence of work, including the detours, so the process is reproducible.

Step	What happened	Why it mattered
1. Load tooling	Loaded the claude-api skill (authoritative guidance for the Anthropic SDK) and the WebFetch tool.	The AI layer calls the Claude API; the skill supplies current model IDs, streaming and refusal-handling rules instead of relying on memory.
2. Name check	Queried pypi.org/pypi/<name>/json for impactlens and datagraph — both returned 404 (available).	A pip-installable library needs a free name; datagraph matched the pitch literally.
3. Architecture	Designed a src-layout package: graph/ (model + graph), extractors/, analysis/, ai/, report.py, cli.py.	Clear separation: deterministic extraction → graph algorithms → AI explanation.
4. Data model	Defined NodeType / EdgeType enums, Node and Edge dataclasses, and the IMPACT_DIRECTION table that says which way change flows for each edge type.	This table is the heart of the library — it makes impact propagation correct without special-casing.

Step	What happened	Why it mattered
5. Graph engine	ImpactGraph over networkx MultiDiGraph: add/merge, resolve references, BFS impact() with depth, impact_tree(), impact_paths(), JSON save/load.	One engine serves every extractor and the CLI.
6. Extractors	Python AST extractor; dbt manifest extractor; sqlglot SQL extractor (optional extra); git diff extractor that maps changed line ranges to functions.	Each emits a graph fragment using shared node-id conventions so fragments merge.
7. Analysis	Deterministic risk scoring (type weights + depth discount + breadth bonus) and a rule-based test recommender; analyze_impact() bundles blast radius + risk + tests + trees.	Produces the opinionated output lineage tools lack.
8. AI layer	explain_impact() using the Anthropic SDK: streaming, claude-opus-5, refusal handling, helpful ImportError when the [ai] extra is missing.	AI explains, never constructs.
9. Report + CLI	rich-based tree/panel renderer; argparse CLI with build / impact / diff / nodes / explain.	Usable from a terminal and from CI.
10. Tests	28 pytest tests with synthetic fixtures (tiny Python project, minimal dbt manifest, SQL files, a throwaway git repo).	Offline, fast (≈ 3 s), covers every extractor and the CLI.
11. Demo	examples/demo: dbt manifest + Python API + one bridging edge; ran datagraph impact models/customer.sql.	Reproduced the pitch output end-to-end; surfaced a resolve() ambiguity bug (fixed).
12. Packaging	python -m build produced wheel + sdist; editable install with [dev] extra.	Confirms pip install works.
13. Open-source prep	.gitignore, GitHub Actions CI, CONTRIBUTING.md, git init + initial commit.	Ready for a public repo.
14. Deep test pass	Fresh venv, install the built wheel, run build/diff/explain on a real git repo written by PowerShell.	Found the UTF-8 BOM bug (extractor silently skipped every file) — fixed in all three extractors with regression tests; hardened CLI explain; added offline AI-layer tests. Test count $\rightarrow$ 34. Second commit.
15. Scenario sweep	Ran impact for ten different targets (four columns, two models, a source, two functions, a file) and a real Python-function edit through datagraph diff.	Proved the engine is generic, not tied to the customer_id example; found that column changes reported nothing (fixed: column-level propagation + rename heuristic) and that legacy Windows consoles crashed on glyphs (fixed: ASCII fallback). Test count $\rightarrow$ 40.

3.2 The v0.3 build — lineage, schema relationships, and “build everything”

v0.3 added the remaining roadmap in one pass, driven by three requests: display lineage; read a schema through a database connection and show all table/column relationships for data analysis; and let an LLM help only where deterministic lineage cannot be derived.

Capability	What was built
Lineage (both directions)	<code>graph.upstream / lineage / upstream_tree</code> ; <code>`datagraph lineage NODE`</code> prints ‘where it comes from’ and ‘what it feeds’ trees; <code>`--html`</code> renders an interactive left-to-right lineage view; <code>`html --all`</code> renders the whole graph.
Schema + connectivity	<code>WarehouseExtractor</code> accepts a DSN (<code>`--warehouse`</code>): <code>information_schema</code> (Postgres, Redshift, Snowflake, BigQuery, DuckDB, MySQL...) or a SQLite file; reads tables, columns+types, FOREIGN KEYS (table↔table and column↔column), view lineage; <code>`datagraph relationships`</code> lists every table’s columns, dependencies and dependents — the schema map for analysts.
LLM fallback for lineage	<code>ai.suggest_lineage</code> uses Claude structured outputs over the schema + unparsed SQL; <code>apply_suggestions</code> adds edges with <code>provenance=llm</code> ; <code>`enrich`</code> / <code>`build --llm-fallback`</code> ; all such edges marked (llm) and excluded by <code>--no-inferred</code> .
Catalog-aware column lineage	<code>`select *`</code> chains (the dbt norm) now resolve through <code>dbt catalog.json / manifest columns / warehouse columns</code> via <code>sqlglot qualify + expand_stars</code> — found by the real <code>jaffle_shop</code> fixture, which initially produced zero column edges.
Code↔data bridge, automatic	SQL embedded in Python/JS/Airflow operators → table edges; table-alias linking joins <code>`schema.table`</code> in code with <code>`db.schema.table`</code> in dbt/warehouse.
Airflow, Lambda, JS/TS, DataHub live	AST-based DAG/task extractor (<code>>></code> , lists, chain, <code>python_callable</code> , SQL in operators); serverless.yml / SAM / CloudFormation extractor (handlers, HTTP APIs, S3/SQS/DynamoDB); regex-based JS/TS extractor; DataHub GraphQL importer (datasets, owners, upstream + fine-grained lineage).
Release plumbing	<code>publish.yml</code> builds, tests, creates a GitHub Release on tags and publishes to PyPI via trusted publishing once enabled; screenshots in docs/images; skill updated with lineage/relationships.

3.1 The v0.2 build — closing the gaps against Graphify and DataHub

After the comparison with Graphify and DataHub, every capability that fits a library (as opposed to a deployed platform) was built directly:

Capability	Borrowed from	What was built
True column-to-column lineage	DataHub / SQLMesh	<code>sqlglot</code> ’s lineage engine traces each output column to its source columns through aliases, CTEs and renames — for SQL files, dbt compiled_code and warehouse view definitions. A renamed column (<code>customer_id</code> →

Capability	Borrowed from	What was built
		customer_key) is now found by derivation, not by name.
OpenLineage import	DataHub / Marquez / Airflow	OpenLineageExtractor reads run events (JSON / NDJSON): datasets, jobs, schema + columnLineage facets, ownership.
DataHub lineage-file import	DataHub	LineageFileExtractor reads DataHub's lineage YAML/JSON (plus simple column lineage and owners).
Warehouse information_schema	DataHub connectors	WarehouseExtractor reads tables, columns (types) and view lineage through any DB-API connection; graph-diff detects schema drift between snapshots.
Ownership / who to notify	DataHub	Owners from dbt meta, exposures and OpenLineage facets; the report prints a Notify list.
Edge provenance	Graphify	Every edge is extracted or inferred; --no-inferred keeps only artifact-backed edges; trees mark heuristics.
Interactive HTML view	Graphify	Self-contained SVG/JS page: layered by depth, click to highlight downstream, search, owners, tests.
Incremental --update, watch, git hook	Graphify	Input fingerprints (outputs excluded), a watch loop, and hook-install for a post-commit rebuild.
paths / hotspots / exports	Graphify	Propagation paths, largest-blast-radius ranking, GraphML / DOT / Cypher / JSON export.
Skill + MCP server	Graphify	skills/datagraph/SKILL.md for Claude Code / Cursor / Codex and datagraph mcp exposing impact, diff, find_nodes, paths, hotspots.
GitHub Action	own roadmap	action.yml posts the blast radius on every PR and can fail the job above a risk level.

Deliberately not built (platform concerns, not library concerns): DataHub's web UI, search/discovery, governance policies and access control, quality monitoring, 50+ live connectors; Graphify's LLM-extracted concepts from docs/PDFs/images and community-detection reports. The honest boundary is documented in the README.

4. Skills, tools and technologies used

4.1 Claude Code skills and tools

Skill / tool	Used for
claude-api skill	Authoritative Anthropic SDK guidance: current model ID (claude-opus-5), default to streaming with get_final_message(), adaptive thinking defaults, handling stop_reason == "refusal", zero-argument Anthropic() client credential resolution.

Skill / tool	Used for
WebFetch	PyPI JSON API lookups to confirm the package name was free.
Write / Edit	Creating the ~30 source, test, and documentation files; surgical fixes (resolve() priority, BOM reads, CLI error handling).
PowerShell / Bash	pip install -e, pytest, python -m build, venv creation, git init/commit, the end-to-end CLI runs, Word automation for the PDF export of this document.
docx skill	This document (docx-js generation, then conversion to PDF).
Memory	Saved a project memory note so future sessions know where the library lives and its design rules.

4.2 Python libraries and standards

Technology	Role in datagraph	Core / optional
ast (stdlib)	Parse Python files: functions, classes, imports, call sites, line spans.	core
networkx ≥ 3.0	MultiDiGraph storage, traversal, all_simple_paths for impact_paths().	core
rich ≥ 13	Terminal panel + tree rendering of the impact report.	core
argparse (stdlib)	CLI with five sub-commands.	core
subprocess + git	git diff --name-only and --unified=0 hunk parsing.	core (git on PATH)
sqlglot ≥ 20	Parse CREATE TABLE/VIEW ... AS SELECT and INSERT statements into table lineage and output columns.	[sql] extra
anthropic ≥ 0.40	Claude API client for the explanation layer.	[ai] extra
pytest ≥ 7	Test suite.	[dev] extra
setuptools + build	PEP 621 pyproject.toml, src layout, console_scripts entry point, wheel + sdist.	build
GitHub Actions	CI matrix (Ubuntu + Windows × Python 3.9/3.11/3.13) plus a build job that uploads dist/.	CI

5. Architecture

5.1 Pipeline

```

Python AST + dbt manifest.json + SQL (sqlglot) + git diff
          | (deterministic extractors)
          v
        Unified ImpactGraph      <- networkx MultiDiGraph, JSON on disk
          | (graph algorithms: BFS with per-edge direction)
          v
    blast radius . risk score . test plan . impact tree
          | (optional)
          v
        Claude explanation (AI explains, never constructs)

```

5.2 Package layout

Path	Responsibility
src/datagraph/__init__.py	Public API re-exports and __version__.
graph/model.py	NodeType, EdgeType, IMPACT_DIRECTION, Node, Edge dataclasses (+ to/from dict).
graph/graph.py	ImpactGraph: add/merge, find/resolve, impact(), impact_tree(), impact_paths(), save/load.
extractors/base.py	Extractor abstract base class — extract() -> ImpactGraph.
extractors/python_extractor.py	Files, functions, classes, CONTAINS, IMPORTS (resolved to project files), best-effort CALLS.
extractors/dbt_extractor.py	Models/seeds/snapshots, sources, exposures, DEPENDS_ON DAG, WRITES_TO tables, declared columns, SQL-file CONTAINS.
extractors/sql_extractor.py	sqlglot lineage: table/view nodes, DEPENDS_ON sources, output COLUMN nodes.
extractors/git_extractor.py	collect_changes() (files + line ranges) and changed_node_ids() (map to file/function/model nodes).
analysis/impact.py	analyze_impact() and the ImpactAnalysis dataclass.
analysis/risk.py	Deterministic risk scoring.
analysis/tests_recommender.py	Rule-based test plan.
ai/explain.py	explain_impact() via the Claude API.
report.py	rich rendering.
cli.py	datagraph build / impact / diff / nodes / explain.
tests/	108 pytest tests across 27 modules, incl. the real jaffle_shop fixture.

Path	Responsibility
examples/demo/	Manifest + Python API + build_demo.py reproducing the pitch.
.github/workflows/ci.yml	CI matrix and build artifact.

5.3 Data model

Node types: file, module, function, class, dbt_model, dbt_source, dbt_seed, dbt_snapshot, exposure, table, view, column, lambda, api, dag, report, dashboard.

Edge types are semantic, and each declares the direction in which change propagates. This is what lets a single BFS give correct answers across code and data:

Edge type	Semantic direction (src → dst)	Impact flows	Meaning
contains	file → function / model → column	forward	changing a file affects everything defined in it
calls	caller → callee	reverse	changing the callee affects its callers
imports	importer → imported	reverse	changing a module affects importers
depends_on	downstream → upstream	reverse	changing upstream data affects dependents
writes_to	model → table	forward	changing a model affects its materialized table
exposes	model → dashboard	forward	changing a model affects the dashboard

5.4 Node id conventions

Every extractor uses the same id scheme, so fragments from different sources merge into one graph and a git-changed path lines up with a dbt model defined in that file:

file:models/customer.sql	func:src/api.py::customers_endpoint
class:src/models.py::Customer	dbt:dim_customer
source:raw.customers	exposure:revenue_report
table:prod.analytics.customer	column:dim_customer.customer_id

5.5 Risk scoring and test recommendations

Risk is deterministic. Each affected node contributes a weight by type — dashboards/reports/exposures 8, APIs 6, lambdas 4, tables/views/dbt models/DAGs 3, seeds/sources 2, functions/classes/columns/files 1 — multiplied by a depth discount (direct hits count fully; deeper hits never below 50 %), plus a breadth bonus for many direct hits (capped at 10). Thresholds: LOW < 6 ≤ MEDIUM < 18 ≤ HIGH < 40 ≤ CRITICAL.

The test recommender maps affected node types to actions: `dbt build --select <model>+` for models, `relationship/not_null` tests for columns, `schema/contract` checks for tables, unit tests for functions, API contract tests, lambda integration tests, manual dashboard validation, DAG test runs.

5.6 The AI layer

`explain_impact(analysis)` serialises the computed `ImpactAnalysis` to JSON and streams a request to Claude (default model `claude-opus-5`, up to 16 000 output tokens) with a system prompt that forbids inventing nodes or dependencies and asks for: what changed, what can break and why (walking the paths), whether the risk level is right, and what to verify before and after deploy. It handles `stop_reason == "refusal"` and raises a clear `ImportError` pointing to `pip install datagraph[ai]` when the SDK is absent. Credentials resolve from `ANTHROPIC_API_KEY` or an ant auth login profile.

6. Worked example — the demo

examples/demo contains a dbt manifest (raw sources → customer → dim_customer → fact_booking → two dashboards) and a small Python service (booking_api.py with fetch_bookings and bookings_endpoint). build_demo.py merges both extractors and adds a single bridging edge:

```
graph.add_edge(Edge(
    src="func:booking_api.py::fetch_bookings",
    dst="table:prod.analytics.fact_booking",
    type=EdgeType.DEPENDS_ON,
))
```

Running datagraph impact models/customer.sql then produced:

```
Change Impact          Changed: models/customer.sql          Risk: HIGH (score 35.0)

models/customer.sql (file)
`-- customer (dbt_model) via contains
    |-- prod.analytics.customer (view) via writes_to
    |-- customer_id / email / created_at (columns) via contains
    `-- dim_customer (dbt_model) via depends_on
        |-- prod.analytics.dim_customer (table) via writes_to
        `-- fact_booking (dbt_model) via depends_on
            |-- prod.analytics.fact_booking (table) via writes_to
            |   |-- fetch_bookings (function) via depends_on
            |   |-- bookings_endpoint (function) via calls
            |-- revenue_report (dashboard) via exposes
            `-- customer_dashboard (dashboard) via exposes

Affected: 8 columns . 3 dbt models . 2 tables . 2 dashboards . 2 functions . 1
view
Recommended tests: dbt build --select customer+ dim_customer+ fact_booking+ ;
schema/contract checks on the 3 relations; unit tests for booking_api.py;
validate both dashboards
```

This is exactly the pitch: a change to one SQL file propagates through the dbt DAG, into warehouse relations, across the code/data bridge into a Python API, and out to the dashboards.

7. How to use it

7.1 Install

```
pip install datagraph          # core: Python + dbt + git extractors
pip install datagraph[sql]     # + SQL lineage via sqlglot
pip install datagraph[ai]     # + Claude explanations
pip install datagraph[all]    # everything
# from the local checkout today:
pip install -e "C:\Users\Hp\Documents\Innovation\datagraph[all]"
```

7.2 Command line

Command	What it does
<code>datagraph build --repo ./src --dbt-manifest target/manifest.json --sql ./sql -o datagraph.json</code>	Run the extractors and save the unified graph.
<code>datagraph impact dbt:customer</code>	Blast radius, risk and test plan for one or more nodes (ids, names or paths). <code>--json</code> for machine output, <code>--max-depth</code> to limit.
<code>datagraph diff --repo . --graph datagraph.json</code>	Blast radius of the current uncommitted change (or <code>--base/--head</code> refs) — maps changed line ranges to the exact functions touched. The CI-friendly command.
<code>datagraph nodes --search customer --type dbt_model</code>	List / search graph nodes.
<code>datagraph explain dbt:customer</code>	Report plus a Claude-written explanation (needs [ai] + credentials).

7.3 Python API

```

from datagraph import ImpactGraph, PythonExtractor, DbtExtractor, analyze_impact

graph = ImpactGraph()
graph.merge(PythonExtractor("./src").extract())
graph.merge(DbtExtractor("target/manifest.json").extract())

analysis = analyze_impact(graph, ["dbt:customer"])
print(analysis.risk)           # {'score': 24.5, 'level': 'HIGH'}
print(analysis.affected)      # {node_id: depth, ...}
print(analysis.recommended_tests)

from datagraph.ai import explain_impact  # optional
print(explain_impact(analysis))

```

8. Testing — what was verified and how

Testing happened in two passes. The first pass was the unit suite written alongside the code; the second was a deliberate “did you test it completely?” pass that installed the built wheel into a clean environment and exercised the real CLI on a real git repository. The second pass is where the most valuable bug was found.

8.1 Automated suite (108 tests, all passing)

Module	Tests	What it pins down
test_python_extractor.py	5	files/functions extracted; import and call edges propagate in reverse; file change hits contained functions; unrelated functions untouched.
test_dbt_extractor.py	8	models/sources/exposures; downstream propagation; SQL file → model; source change hits everything; upstream not affected; materialized tables linked; risk + dbt build recommendation; resolve by name.
test_sql_extractor.py	3	table lineage through two CREATE TABLE AS SELECT files; output columns; downstream-only.
test_graph.py	6	merging code + data graphs with a bridge edge; save/load round-trip; impact tree shape; impact paths; max_depth; JSON-serialisable analysis.
test_git_extractor.py	2	uncommitted change maps to the exact function (and not its neighbour); diff impact reaches the caller.
test_cli.py	4	build + impact --json; unknown node → exit 2; missing graph → exit 2; nodes listing.
test_bom_files.py	3	UTF-8 BOM files parse in all three extractors (regression).
test_ai_layer.py	3	helpful ImportError without anthropic; request payload contains the deterministic analysis and the no-invention system prompt; refusal handled — all against a stub SDK, offline.
test_column_impact.py	4	a column change reaches downstream models and exposures; same-named downstream columns flagged, unrelated siblings not; tree roots at the column; model change still flags its own columns.
test_report.py	2	report renders with glyphs on UTF-8 and falls back to ASCII on a cp1252 (legacy Windows) console without crashing.

8.2 End-to-end verification from the built wheel

1. Built wheel + sdist with `python -m build`.
2. Created a fresh venv, installed the wheel with the [sql] extra, confirmed the `console_scripts` entry point metadata (`datagraph = datagraph.cli:main`) and imports.

3. Created a throwaway git repo with PowerShell-written Python files and the demo manifest; ran datagraph build; edited one function; ran datagraph diff.
4. Result after fixes: the diff command detected the uncommitted change, flagged db.py and load_customers specifically (not load_orders), and propagated to customers_endpoint via the call edge — Risk MEDIUM, exit 0.
5. Verified datagraph explain exits 2 with a clear message (no traceback) when the [ai] extra is absent.

8.3 Bugs found and fixed along the way

Bug	How it was found	Fix
resolve() treated a path shared by a file node and the dbt model defined in that file as ambiguous.	Running the demo CLI.	Resolution order: exact id → exact name → exact path → file node on that path → unique substring.
UTF-8 BOM (default for PowerShell / many Windows editors) made ast.parse fail; the Python extractor silently skipped every file (“0 nodes”).	Wheel-based end-to-end run — unit fixtures write BOM-less files so pytest never saw it.	All extractors read utf-8-sig; three regression tests added.
datagraph explain dumped a raw traceback when anthropic was missing.	Same end-to-end pass.	Catch ImportError, print guidance, exit 2.
Stale wheel: an e2e run used a wheel built before a fix.	Process slip during testing.	Rebuild before every reinstall; noted as a lesson.
A column change reported nothing affected (column nodes only had an incoming contains edge).	Scenario sweep: impact column:fact_booking.amount returned an empty blast radius.	Column changes now propagate through the owning model/table to everything downstream and flag same-named downstream columns (rename heuristic).
UnicodeEncodeError on legacy Windows consoles (cp1252) — the report’s ⚠ / 🔴 / ✓ glyphs crashed rich.	Running datagraph diff from Git Bash.	ASCII fallback icons and markers when the console is not UTF-8.
rich interpreted the ASCII icon [dbt] as markup and stripped it.	The new cp1252 rendering test.	Escape icons and node names with rich.markup.escape.

8.4 Honest remaining gaps

- **Live Claude API call:** not executed — no ANTHROPIC_API_KEY on the build machine. Everything around the call is pinned by stub tests; one run of datagraph explain with a key closes it.
- **Python 3.9–3.13:** the build machine only has 3.14; the CI matrix covers the older versions on first push.

- **Console .exe launcher:** this machine's security policy blocks every pip-generated .exe (environmental, not a packaging defect); verified via entry-point metadata and `python -m datagraph.cli`.

9. How datagraph differs from what already exists

Parts of this problem are well covered by existing tools. The honest positioning is that datagraph is a complement to the lineage platforms, not a replacement — the fast, local, PR-time layer that also sees your code.

Tool / category	What it does well	Where it stops	datagraph's difference
DataHub, OpenLineage, Marquez	Enterprise data lineage across warehouses, pipelines, BI; runtime-collected; DataHub has an impact view.	Data plane only — no application code; a platform to deploy and ingest into, not a 5-second PR check.	Adds the code plane and the diff-to-function entry point; runs locally from pip in CI.
dbt (dbt ls --select model+)	Model-level lineage and exposures inside one dbt project.	Ends at the dbt boundary; no Python callers, no git-diff mapping, no risk or test plan.	Merges dbt lineage with code, scores risk, recommends tests.
Graphify, Code-Graph-RAG, Sourcegraph-style	Build graphs from source code, usually to feed an LLM for code Q&A.	Code plane only; the LLM often participates in interpreting the graph.	Deterministic graph; crosses into tables and dashboards.
sqllineage, sqlglot, SQLMesh	SQL parsing and (for SQLMesh/DataHub) real column-level lineage.	A layer, not a PR-time impact tool.	datagraph uses sqlglot for this layer; adds everything above it.

9.1 Side by side: Graphify · DataHub · OpenLineage · datagraph

Graphify helps an AI assistant understand a repo; DataHub is the company-wide catalog; OpenLineage is the open standard lineage events travel in; datagraph is the pre-merge check and lineage/relationship explorer that also reads your code — and imports the other two rather than competing with them.

	Graphify	DataHub	OpenLineage	datagraph
What it is	Skill: folder → knowledge graph for assistants	Deployed metadata platform / catalog	Open standard + spec for lineage events (Marquez reference server)	pip library + CLI + skill
Question answered	Understand this repo	What exists, who owns it, is it healthy?	What did this job read/write at run time?	What breaks if I change this? Where does it come from? How are tables related?
Inputs	13 languages, docs, PDFs, images	50+ connectors, OL events	Emitters (Airflow, Spark, dbt, Flink...)	Python/JS AST, dbt manifest+catalog, SQL, warehouse information_schema (FKs), git diff, Airflow,

	Graphify	DataHub	OpenLineage	datagraph
				Lambda, OpenLineage, DataHub
Graph built by	AST + Claude for non-code	Ingestion connectors	The emitting jobs	Deterministic extractors; optional llm fallback, tagged
Knows application code	Yes (structure)	No	No	Yes — functions, calls, SQL-in-code, handlers, callables
Column lineage	No	Yes	Yes (facet)	Yes (sqlglot, catalog-aware; imports OL/DataHub)
Foreign keys / schema map	No	Yes	No	Yes (relationships, FK edges, drift)
Impact direction + risk + tests	No	Impact view	No	Yes, across code and data
Entry point	A folder	A dataset	A job run	A git diff, or any node
Display	HTML graph, wiki	Web UI	Via a backend	Interactive HTML (impact/lineage/whole), terminal trees, GraphML/DOT/Cypher
Infrastructure	None	Platform	Needs a backend	None — pip, JSON, CI; skill + MCP
AI role	Extracts concepts	n/a	n/a	Explains; optional tagged suggestions

9.2 What Graphify has that datagraph does not

- 13 languages via tree-sitter (we parse Python only).
- Docs / Markdown / PDFs / images turned into concepts by Claude — deliberately excluded: our rule is no LLM in graph construction.
- Community detection, god-node report, interactive graph.html, Obsidian vault, wiki generation.
- Incremental --update, --watch, SHA256 cache, post-commit hook; query / path / explain commands; SVG, GraphML and Neo4j exports.
- Edge provenance tags (EXTRACTED / INFERRED / AMBIGUOUS) — worth adopting for our name-resolved call edges and same-name column matches.
- A /graphify skill and MCP server inside Claude Code, Cursor, Codex, Gemini CLI — the distribution channel that made it famous, and the biggest gap on our side.

9.3 What DataHub has that datagraph does not

- Ingestion from 50+ systems (warehouses, orchestrators, BI tools, quality tools) and runtime-collected lineage; true table- and column-level lineage through BI dashboards.
- Search and discovery, dataset profiles and statistics, an impact-analysis UI.
- Governance: ownership, glossary, tags, domains, data products, PII classification, deprecation; access policies, roles, audit.
- Data quality assertions, incidents, freshness/volume SLAs, data contracts, schema-change detection against live metadata.
- A platform: GraphQL/REST APIs, Kafka metadata stream, Actions framework, web UI, multi-tenant deployment.

9.4 What datagraph has that neither does

1. **One graph spanning both worlds.** A git diff on a Python file and a dbt model change land in the same graph, so impact propagates code → table → dbt → dashboard and back. Graphify stops at code structure; DataHub stops at the data plane.
2. **Diff-to-function granularity as the entry point.** The unit of analysis is the change you are about to merge; changed line ranges are mapped to the exact functions touched.
3. **A verdict, not a map.** Risk level and a concrete test plan, computed deterministically.
4. **Zero infrastructure, offline, deterministic.** pip install, a JSON graph, one command in CI; nothing leaves the repo; the model only narrates an analysis it cannot alter.

9.5 What datagraph deliberately does not try to be

A catalog (search, glossary, ownership, policies, quality monitoring — use DataHub) or a general code-understanding graph with docs/PDF/image ingestion (use Graphify). The roadmap item that ties the three together is an OpenLineage / DataHub import so datagraph inherits enterprise lineage and adds the code side and the PR-diff entry point on top; and a /datagraph skill + MCP server so it reaches the same audience Graphify did.

9.6 Where existing tools are ahead (known gaps)

- **Coverage:** lineage platforms see runtime lineage across dozens of systems; datagraph now imports that lineage (OpenLineage, DataHub files, information_schema) but still parses only Python for code. Python call resolution is name-based (tagged_inferred), and code↔data bridge edges are declared by hand — auto-detecting them is the highest-value roadmap item.
- **Column-level lineage:** now real (sqlglot-derived) wherever SQL is available — compiled dbt code, SQL files, warehouse view definitions, OpenLineage columnLineage facets. Without SQL the same-name heuristic remains as a marked-inferred fallback.
- **Freshness:** a stored graph goes stale; mitigated by rebuilding in CI, but lineage platforms update continuously.

10. Packaging and open-source readiness

10.1 Packaging

- PEP 621 pyproject.toml, setuptools backend, src layout, requires-python ≥ 3.9 , MIT license, classifiers and keywords set.
- Core dependencies kept light (networkx, rich); heavy ones behind extras: [sql] sqlglot, [ai] anthropic, [all], [dev] pytest.
- console_scripts entry point: datagraph = datagraph.cli:main.
- dist/ contains datagraph-0.1.0-py3-none-any.whl and datagraph-0.1.0.tar.gz; publishing is `python -m twine upload dist/*`.

10.2 Repository hygiene

- README.md — pitch, install, CLI and Python quick starts, propagation rules, id conventions, roadmap.
- CONTRIBUTING.md — setup, the three project principles, a recipe for adding an extractor, PR rules.
- LICENSE (MIT), .gitignore, .github/workflows/ci.yml (test matrix + build artifact).
- Git history: a452a29 initial release; ced063a BOM fix + AI-layer tests; 7f9b303 documentation; 29c0378 column-level impact + ASCII console fallback; 3b1d1e9 markup escaping.

10.3 Publishing to GitHub (next step)

The GitHub CLI is not installed on the build machine, so the repo has not yet been pushed. Either:

1. Install GitHub CLI (`winget install --id GitHub.cli -e`), run `gh auth login`, then `gh repo create datagraph --public --source . --push`; or
2. Create an empty public repo on github.com, then `git remote add origin https://github.com/<you>/datagraph.git` and `git push -u origin main` (a browser sign-in appears on first push).

Before announcing: replace the placeholder Homepage URL in pyproject.toml and <owner> in CONTRIBUTING.md with the real GitHub username, and consider adding the demo output as a screenshot near the top of the README.

11. Roadmap

Item	Value	Effort
/datagraph skill + MCP server for Claude Code / Cursor / Codex	Distribution — the channel that made Graphify famous	Low
Interactive HTML blast-radius view; incremental --update / --watch; edge provenance tags	Shareable output, freshness, honesty about inferred edges	Medium
Auto-detect code $\leftrightarrow$ data bridge edges (SQL strings in Python, warehouse connector / boto3 calls)	Highest — removes the one manual step	Medium
Column-level lineage through dbt compiled SQL	High for rename / type-change impact	Medium–high
OpenLineage / DataHub import	Inherit enterprise lineage; position as complement	Medium

Item	Value	Effort
Airflow DAG and AWS Lambda extractors	Completes the pitched chain	Medium
GitHub Action posting the impact report on every PR	Turns users into daily users	Low
Snowflake metadata extractor (information_schema)	Real tables/columns without a manifest	Medium

12. v0.5 — data profiling, knowledge base for AI assistants, extractor plugins

After v0.4 the question was: can all the ‘data things’ live in this one library, and can the code + data graph be turned into a knowledge graph that AI assistants use? v0.5 answers yes on both counts without changing the core principle — everything below is generated deterministically from the graph; the LLM (if any) only reads it.

12.1 What was added

Capability	Command / API	What it does
Data profiling	<code>datagraph profile --warehouse DSN · profile_warehouse(conn, graph)</code>	Through any DB-API / SQLAlchemy / SQLite connection: row count, freshness (max of date-like columns), per-column null %, distinct count, min/max and top values (sampled). Stored on the node as <code>meta['profile']</code> ; shown in relationships, lineage HTML tooltips and context packs.
Profile-aware risk	<code>automatic</code>	Empty tables count half, tables with $\geq 1M$ rows count 1.5× in the risk score; distinct counts are passed to the LLM lineage fallback so it can spot join keys.
Context pack	<code>datagraph context NODE · context(graph, id) · MCP tool context</code>	One compact text block per node: description, owner, columns (+type, pk, profile, where each column comes from), upstream, downstream, table relationships, dbt tests, heuristic-edge warning, risk if changed + test plan, and the SQL that builds it.
Knowledge base (wiki)	<code>datagraph wiki -o DIR · build_wiki(graph, dir)</code>	Markdown knowledge base: <code>index.md</code> grouped by type, one page per table / model / source / function / DAG / task / lambda / API / dashboard with cross-links, <code>llms.txt</code> for RAG tools, and <code>GRAPH_REPORT.md</code> .
Graph report	<code>part of wiki</code>	Hotspots (largest blast radius), high-impact dbt models without tests, nodes without an owner, roots (raw inputs), leaves (nothing

Capability	Command / API	What it does
		downstream — deletion candidates), count of heuristic edges.
dbt enrichment	DbtExtractor	Each model now carries its compiled SQL (meta['sql']) and the names of the dbt tests that cover it (meta['tests']).
Extractor plugins	register(ExtractorPlugin(...)) · entry point datagraph.extractors · datagraph plugins	Any package can add an extractor (Looker, Tableau, Dagster, Kafka...). It automatically becomes a datagraph build --<name> flag with its own --<name>-<option> flags and merges into the same graph.

12.2 Why this matters

- **One graph, many consumers.** The same deterministic graph now serves PR checks (impact), data analysis (lineage, relationships, profiles) and AI assistants (context, wiki, MCP). Nothing is built twice.
- **Assistants explain instead of guessing.** A context pack is what Claude Code / Cursor / a RAG bot needs before touching a table or model; because it is generated from artifacts, hallucinated lineage is replaced by extracted lineage with provenance tags.
- **Extensible without forking.** Plugins let teams cover their own BI / orchestration tools; the registry is tested with a toy BI extractor that appears in impact results through the shared table node.

12.3 Tests added

File	Tests	What is verified
test_profiling.py	3	SQLite fixture: row counts, 10% nulls, distinct counts, min/max, freshness; profile flows into schema_summary for the LLM and into risk; CLI profile --tables and relationships --json show profiles.
test_knowledge.py	6	context pack content (upstream, downstream, risk, no-match), SQL and tests captured from a manifest, wiki files (index, GRAPH_REPORT, llms.txt, node pages with links), CLI wiki/context, MCP context tool.
test_plugins.py	3	Programmatic registration becomes a build flag with options; plugin-only build allowed; non-graph return raises TypeError.

Total: 120 tests, all offline, ~15 s. Released as v0.5.0 (tag + GitHub Release).

Appendix A — Command reference

```
# development
pip install -e ".[dev]"          # install with tests + sqlglot
python -m pytest -q              # 120 tests
python -m build                  # wheel + sdist into dist/

# demo
cd examples/demo && python build_demo.py
datagraph impact models/customer.sql --graph datagraph.json

# everyday
datagraph build --repo ./src --dbt-manifest target/manifest.json -o
datagraph.json
datagraph diff --repo . --graph datagraph.json
datagraph impact dbt:customer --json
datagraph nodes --search booking
datagraph explain dbt:customer --model claude-opus-5

# v0.5 — knowledge base for AI assistants
datagraph profile --warehouse prod.db          # row counts, freshness, null %,
distinct, min/max, top values
datagraph context dim_customer                # compact text pack for one node
datagraph wiki -o kb/                         # index.md, GRAPH_REPORT.md,
llms.txt, nodes/*.md
datagraph plugins                             # installed extractor plugins
datagraph mcp --graph datagraph.json          # tools incl. context
```

Appendix B — File inventory

Source (src/datagraph): `__init__.py`, `cli.py`, `report.py`, `graph/{__init__,model,graph}.py`,
`extractors/{__init__,base,python_extractor,dbt_extractor,sql_extractor,git_extractor}.py`,
`analysis/{__init__,impact,risk,tests_recommender}.py`, `ai/{__init__,explain}.py`.

Tests: `conftest.py`, `test_python_extractor.py`, `test_dbt_extractor.py`, `test_sql_extractor.py`,
`test_graph.py`, `test_git_extractor.py`, `test_cli.py`, `test_bom_files.py`, `test_ai_layer.py`.

Project: `pyproject.toml`, `README.md`, `CONTRIBUTING.md`, `LICENSE`, `.gitignore`,
`.github/workflows/ci.yml`, `examples/demo/{manifest.json, app/booking_api.py, build_demo.py}`,
`docs/` (this document).

Appendix C — Design decisions worth remembering

- **Semantic edges + a direction table** instead of storing edges in impact direction: keeps extractors natural (“model depends on source”) and the traversal correct.
- **Placeholder nodes on dangling edges** (a dbt model may depend on a node that a later extractor defines): merging never loses an edge.
- **Name-based call resolution** is deliberately best-effort; false positives are preferable to missed impact for a safety tool, and the tree shows “via calls” so reviewers can judge.
- **No LLM in the graph** — tests assert the exact prompt payload so the AI layer can never quietly become load-bearing.